

# Design Notes for Journal CLI

Manjunatha Sarma Majety

January 5, 2026

## Abstract

This document serves as a personal log and design record for the `jrn1` CLI tool. It captures the evolution of the project, lessons learned, and insights into the thought process behind each design choice.

## Contents

<b>1</b>	<b>Motivation</b>	<b>2</b>
<b>2</b>	<b>Design Evolution</b>	<b>2</b>
2.1	Phase 1: Structure Before Clarity . . . . .	2
2.2	Phase 2: OOP Refactor . . . . .	2
2.3	Post-Refactor Milestones: Atomic Saves and Configuration . . . . .	3
2.4	The Next Big Thing: Date-Based Filters & Expanding the Config File . . . . .	4
2.5	Adding Backups and the Path to Publishing . . . . .	4
2.6	A Quiet, Strange Kind of Ending . . . . .	5
2.7	Local Journals . . . . .	5
2.8	Reflections and the End . . . . .	5
2.9	Work in Progress and Future Plans . . . . .	6

# 1 Motivation

The goal of this project was to create a simple, fast, and minimalistic command-line journaling tool that felt native to Unix-style workflows. I wanted something I could rely on daily—lightweight, intuitive, and built on a design philosophy centered around clarity and control.

## 2 Design Evolution

### 2.1 Phase 1: Structure Before Clarity

Initially, the entire program was written in a purely procedural style. It consisted of two main functions—`display()` and `addEntry()`—which handled reading and writing entries respectively. A tiny helper function took care of timestamp generation and UNIX-to-human time conversion.

While this design gave the program a basic skeleton, it was fragile. A single malformed line in the journal file could break `getline()`, and adding new features often meant rewriting large portions of the code. It worked, technically, but it felt far too brittle for something I wanted to trust long-term.

Still, this phase mattered. It was my first fully coherent, working program—something I could actually use. I picked up lessons about `CMake`, time libraries (in a very brute-force kind of way), ANSI color codes, and the importance of clean header/implementation separation. Those small lessons ended up laying the groundwork for everything that came next.

### 2.2 Phase 2: OOP Refactor

This was the first major transformation: rebuilding the tool from a procedural script into an object-oriented system. Classes felt intimidating at first, but once they started connecting in predictable ways, the structure of the program became clearer, more modular, and much easier to extend.

During this refactor, I reworked the display logic to accept a wider range of arguments and specifiers. One unexpectedly tricky problem was parsing expressions like "`*3`" or "`3*`". My first attempt was a naive ASCII-based trick (`num - '48'`), which broke instantly for multi-digit numbers. I eventually replaced it with `stoi()`, which forced me to actually learn how string parsing and type conversion work in practice.

This phase gave me a deeper sense of how C++ behaves beneath its surface. Not just what the language can express, but how its moving parts fit together. It felt like progress in both code and mindset.

## 2.3 Post-Refactor Milestones: Atomic Saves and Configuration

After the OOP refactor, the project crossed its biggest stability milestone so far: **atomic file saves** and **robust configuration handling**. These changes fundamentally altered how safe and trustworthy the tool became.

Originally, the save logic was dangerously naive. The journal was cleared as soon as it was read into memory, and then rewritten completely. A crash, a power loss, or even an interrupted write meant the entire journal could evaporate. At the time, I underestimated how catastrophic that was.

Eventually, I realised that correctness is meaningless without *durability*. So I rewrote the entire save pipeline around POSIX guarantees:

- Write all entries to a temporary `.tmp` file.
- Flush it explicitly to disk via `fsync()`.
- Close the file descriptor.
- Atomically replace the original using `rename()`.
- `fsync()` the parent directory for extra safety.

This ensured that the journal could never exist in a partially-written state. Either the new version replaces the old one entirely, or the old version remains. There is no corrupted in-between. This alone elevated the project from “toy” to “tool I can trust.”

Around this period, I also implemented full configuration handling. The program now respects XDG base directories when available, falling back to `$HOME` otherwise. If no config exists, one is generated automatically with sane defaults. Directories and journal files are created lazily and safely.

This phase also pushed me into a long-overdue parser refactor. The earlier argument-handling logic had grown into a messy tangle of conditionals, tightly coupled to execution. The new dispatcher-based design—inspired loosely by `git`—cleanly separates *what* the user wants from *how* the program executes it. No new features appeared, but the internal clarity improved dramatically.

I also had to face an uncomfortable truth: it is almost impossible to write fully error-proof code from the beginning. Defensive checks should have been added incrementally, not as a giant cleanup pass. But tackling it all taught me about:

- POSIX file descriptors and low-level I/O
- Command dispatchers and their simplicity
- Atomic filesystem guarantees and their limits
- The difference between logical correctness and durability
- Why safety mechanisms must be designed, not bolted on later

More than any individual feature, this milestone reshaped how I think about reliability. The journal no longer merely *works*—it now fails *responsibly*.

## 2.4 The Next Big Thing: Date-Based Filters & Expanding the Config File

In all honesty, every time I thought the tool was finally approaching a neat, sensible conclusion—something that would let me breathe for a moment—it managed to grow again. The next target was proper time-based filtering using `--before` and `--after`. On paper, the logic was simple: use `std::lower_bound` and `std::upper_bound` to mark the search boundaries. Reality, of course, had ahem... opinions.

The real challenge was composability. The range-based filters and time-based filters each worked fine alone. But making them coexist without stepping on each other's toes? That took some thinking.

My first attempt was a quick, almost improvised logic. Surprisingly, it worked—for a handful of cases. But it broke on combinations like `--before` with `3*`, or `--after` with `*3`.

Eventually, early one morning, the actual solution arrived through a tiny bit of math. The two filter types didn't have equal priority the way I assumed. Time-based filters had to dominate, with range-based ones slicing *within* those constraints. Once I rewired the logic around that idea, everything fell neatly into place.

Math saves the day again.

Around this time, I also expanded the configuration system. What used to be a simple path variable grew into something bigger: backup paths, default color codes, and more flexible initialization logic. I split the old single-purpose constructor into clearer methods, making the design more modular and future-proof.

This stage felt like opening a door to a room I hadn't noticed before. Suddenly, all the tiny choices I'd postponed started demanding attention. But fixing them wasn't a burden—it felt like the project was maturing, piece by piece, into something I could be genuinely proud of.

## 2.5 Adding Backups and the Path to Publishing

To be very honest, backups were the only addition that felt simple compared to the work I had done over the past month. They were implemented as a slightly modified version of the atomic write code I had already written for the save function. All things considered, what remained was adding a neat README, man pages, a PKGBUILD, and finally publishing the tool to the AUR—tasks that were accomplished soon enough.

## 2.6 A Quiet, Strange Kind of Ending

When the color configuration finally clicked into place, I felt an unexpected wave of relief wash over me. Not because the project was truly finished—it isn’t, and I still have edge cases and error paths scattered around like crumbs—but because it finally felt *complete enough*. The tool had taken on a stable, coherent shape: a shape I could trust, and one I could step away from without worry.

It was a strange feeling, almost bittersweet. I watched this little program grow from a clumsy procedural script into something deliberate, reliable, and strangely alive. And, in a quiet way, it watched me grow too—through mistakes, rewrites, midnight bug hunts, and those tiny breakthroughs that feel bigger than they are. Now it feels as though we are about to part ways. I will still use it often, of course; that is why I built it in the first place. But there is a soft, lingering melancholy in knowing that this chapter is closing, and that the tool and I are no longer evolving together at the same pace.

## 2.7 Local Journals

After building this humongous tool for my day-to-day use, I ironically ran into another need: local journals that could act as notebooks for all the different projects I was working on. The exciting part was that it didn’t require a full redesign—just a few thoughtful additions to the existing framework.

This led to a mini-refactor, introducing `init` and the `--local`/`--global` flags. The design is simple: unless `--global` is explicitly specified, if both global and local journals exist, any additions, `show`, or other operations default to the local journal.

Alongside this, I implemented a small but convenient feature: the previous entry’s tag is used as the default for the next one. In project notebooks or other scenarios where the tag isn’t simply `jrn1`, this removes the need to repeatedly pass `--tag`—once is enough at initialization. These few adjustments made the tool feel significantly more flexible and usable across a broader set of workflows.

## 2.8 Reflections and the End

Once I published the tool to the AUR and shared a post on Reddit, I got some thoughtful feedback. One user pointed out that using `*5` range notation could conflict with shell globs and suggested that regex might be more intuitive. I wasn’t fully convinced, since my experience suggested users would adapt to familiar conventions. Instead, I adopted Python-style slicing notation `(:5)` alongside the new `--first` and `--last` flags for clarity and predictability.

By this stage, I noticed something deeper: the way I write code had fundamentally changed. Writing modular code had become almost second nature, a kind of muscle memory. Anticipating edge cases was no longer exhausting—adding new features, which at the start of the project had felt like massive refactors,

now felt like manageable side quests. It was a subtle, satisfying shift: the tool was stable, extensible, and my own development style had matured alongside it.

## 2.9 Work in Progress and Future Plans

In my opinion, this tool has reached a level of functionality and correctness that exceeds what I initially needed—or even expected—for my day-to-day use. That said, there are still features I hope to add in the future. For example, exporting journal entries to `JSON` or `Markdown` format is on the roadmap, though not in the near term.

Beyond that, I remain open to suggestions and ideas for new features or improvements.